

## ECE/CS 250 Computer Architecture

Multicore and Multithreaded Processors  
Copyright Daniel J. Sorin  
Duke University

## Multicore and Multithreaded Processors

---

- Why multicore?
- Thread-level parallelism
- Multithreaded cores
- Multiprocessors
- Design issues
- Examples

## Readings

---

- Patterson and Hennessy
  - Chapter 6
  - At least one recent research paper!

## Why Multicore?

---

- Why is everything now multicore?
  - This is a fairly new trend
- Reason #1: Running out of “ILP” that we can exploit
  - Can't get much better performance out of a single core that's running a single program at a time
- Reason #2: Power/thermal constraints
  - Even if we wanted to just build fancier single cores at higher clock speeds, we'd run into power and thermal obstacles
- Reason #3: Moore's Law
  - Lots of transistors → what else are we going to do with them?
  - Historically: use transistors to make more complicated cores with bigger and bigger caches
  - But this strategy has run into problems

## How do we keep multicores busy?

- Single core processors exploit ILP
- Multicore processors exploit **TLP: thread-level parallelism**
- What's a thread?
  - A program can have 1 or more threads of control
  - Each thread has own PC
  - All threads in a given program share resources (e.g., memory)
- OK, so where do we find more than one thread?
- Option #1: Multiprogrammed workloads
  - Run multiple single-threaded programs at same time
- Option #2: Explicitly multithreaded programs
  - Create a single program that has multiple threads that work together to solve a problem

## Parallel Programming

- How do we break up a problem into sub-problems that can be worked on by separate threads?
- ICQ: How would you create a multithreaded program that searches for an item in an array?
- ICQ: How would you create a multithreaded program that sorts a heap?
- Fundamental challenges
  - Breaking up the problem into many reasonably sized tasks
    - What if tasks are too small? Too big? Too few?
  - Minimizing the communication between threads
    - Why?

## Writing a Parallel Program

- Compiler can turn sequential code into parallel code
  - Just as soon as the Cleveland Indians win the World Series
- Can use an explicitly parallel language or extensions to an existing language
  - Map/reduce (Google), Hadoop
  - Pthreads
  - Java threads
  - Message passing interface (MPI)
  - CUDA
  - OpenCL
  - High performance Fortran (HPF)
  - Etc.

## Parallel Program Challenges

- Parallel programming is HARD!
  - Why?
- Problem: #cores is increasing, but parallel programming isn't getting easier → how are we going to use all of these cores???

## HPF Example

```
forall(i=1:100, j=1:200){  
    MyArray[i,j] = X[i-1, j] + X[i+1, j];  
}
```

// "forall" means we can do all i,j combinations in parallel  
// I.e., no dependences between these operations

## Some Problems Are "Easy" to Parallelize

- Database management system (DBMS)
- Web search (Google)
- Graphics
- Some scientific workloads (why?)
- Others??

## Multicore and Multithreaded Processors

- Why multicore?
- Thread-level parallelism
- [Multithreaded cores](#)
- Multiprocessors
- Design issues
- Examples

## Multithreaded Cores

- So far, our core executes one thread at a time
- Multithreaded core: execute multiple threads at a time
- Old idea ... but made a big comeback fairly recently
- How do we execute multiple threads on same core?
  - Coarse-grain switching
  - Fine-grain switching
  - Simultaneous multithreading (SMT) → "hyperthreading" (Intel)
- Benefits?
  - Better instruction throughput
    - Greater resource utilization
    - Tolerates long latency events (e.g., cache misses)
  - Cheaper than multiple complete cores

## Multiprocessors

- Multiprocessors have been around a long time ... just not on a single chip
  - Mainframes and servers with 2-64 processors
  - Supercomputers with 100s or 1000s of processors
- Now, multiprocessor on a single chip
  - "multicore processor" (sometimes "chip multiprocessor")
- Why does "single chip" matter so much?
  - ICQ: What's fundamentally different about having a multiprocessor that fits on one chip vs. on multiple chips?

## Multicore and Multithreaded Processors

- Why multicore?
- Thread-level parallelism
- Multithreaded cores
- Multiprocessors
- Design issues
- Examples

## Multiprocessor Microarchitecture

- Many design issues unique to multiprocessors
  - Interconnection network
  - Communication between cores
  - Memory system design
  - Others?

## Interconnection Networks

- Networks have many design aspects
  - We focus on one design aspect here (topology) → see ECE 552 (CS 550) and ECE 652 (CS 650) for more on this
- Topology is the structure of the interconnect
  - Geometric property → topology has nice mathematical properties
- Direct vs Indirect Networks
  - Direct: All switches attached to host nodes (e.g., mesh)
  - Indirect: Many switches not attached to host nodes (e.g., tree)

## Direct Topologies: k-ary d-cubes


- Often called k-ary **n**-cubes
- General class of regular, **direct** topologies
  - Subsumes rings, tori, cubes, etc.
- **d** dimensions
  - 1 for ring
  - 2 for mesh or torus
  - 3 for cube
  - Can choose arbitrarily large d, except for cost of switches
- **k** switches in each dimension
  - Note: k can be different in each dimension (e.g., 2,3,4-ary 3-cube)

## Examples of k-ary d-cubes (for N cores)

- 1D Ring = k-ary 1-cube
  - $d = 1$  [always]
  - $k = N$  [always] = 4 [here]
  - Ave dist = ?



- 2D Torus = k-ary 2-cube
  - $d = 2$  [always]
  - $k = \log_d N$  (always) = 3 [here]
  - Ave dist = ?



## k-ary d-cubes in Real World


- Compaq Alpha 21364 (and 21464, R.I.P.)
  - 2D torus (k-ary 2-cube)
- Cray T3D and T3E
  - 3D torus (k-ary, 3-cube)
- Intel's MIC (formerly known as Larrabee)
  - 1D ring
- Intel's SandyBridge (one flavor of core i7)
  - 2D mesh

## Indirect Topologies

- Indirect topology – most switches not attached to nodes
- Some common indirect topologies
  - Crossbar
  - Tree
  - Butterfly
- Each of the above topologies comes in many flavors

## Indirect Topologies: Crossbar

- Crossbar = single switch that directly connects  $n$  inputs to  $m$  outputs
  - Logically equivalent to  $m \times n$  muxes
- Very useful component that is used frequently



© Daniel J. Sorin

ECE 250

## Indirect Topologies: Trees


- Indirect topology – most switches not attached to nodes
- **Tree**: send message up from leaf to closest common ancestor, then down to recipient
- $N$  host nodes at leaves
- $k$  = branching factor of tree ( $k=2 \rightarrow$  binary tree)
- $d$  = height of tree =  $\log_k N$

© Daniel J. Sorin

ECE 250

## Indirect Topologies: Fat Trees

- Problem with trees: too much contention at or near root
- **Fat tree**: same as tree, but with more bandwidth near the root (by adding multiple roots and high order switches)




© Daniel J. Sorin

ECE 250

## Indirect Topologies: Butterflies

- Multistage: nodes at ends, switches in middle
- Exactly one path between each pair of nodes
- Each node sees a tree rooted at itself



© Daniel J. Sorin

ECE 250

## Indirect Topologies: More Butterflies

- In general, called **k-ary, n-flies**
  - **n stages of radix-k switches**
- Have many nice features, esp.  $\log_n$  distances
- But conflicts cause **tree saturation**
  - How can we spread the traffic more evenly?

### Benes (pronounced “BEN-ish”) Network



- Routes all permutations w/o conflict
- Notice similarity to fat tree (fold in half)
- **Randomization is major breakthrough**

## Indirect Networks in Real World (ancient)

- Thinking Machines CM-5 (really old machine)
  - Fat tree
- Sun UltraEnterprise E10000 (old machine)
  - 4 trees (interleaved by address)
- And lots and lots of buses!

## Multiprocessor Microarchitecture

- Many design issues unique to multiprocessors
  - Interconnection network
  - **Communication between cores**
  - Memory system design
  - Others?

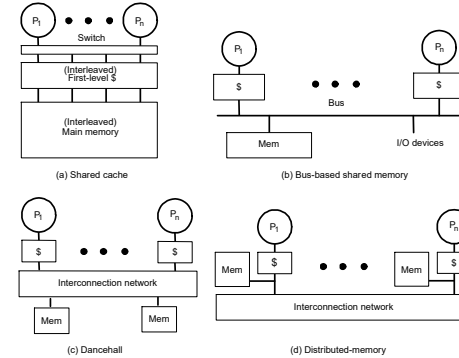
## Communication Between Cores (Threads)

- How should threads communicate with each other?
- Two popular options
- **Shared memory**
  - Perform loads and stores to shared addresses
  - Requires synchronization (can't read before write)
- **Message passing**
  - Send messages between threads (cores)
  - No shared address space


## What is (Hardware) Shared Memory?

- Take multiple microprocessors
- Implement a memory system with a single global physical address space (usually)
  - Special HW does the "magic" of cache coherence

## Some (Old) Memory System Options



## A (Newer) Memory System Option




## Cache Coherence

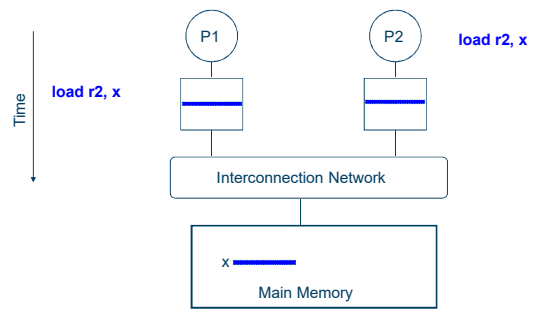
- According to Webster's dictionary ...
  - **Cache**: a secure place of storage
  - **Coherent**: logically consistent
- Cache Coherence: keep storage logically consistent
  - Coherence requires enforcement of 2 properties per block
    - 1) At any time, only one writer or  $\geq 0$  readers of block
      - Can't have writer at same time as other reader or writer
    - 2) Data propagates correctly
      - A request for a block gets the most recent value



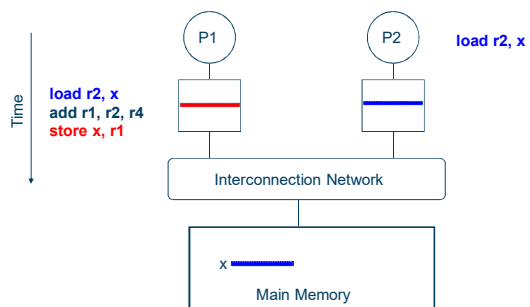
### Cache Coherence Problem (Step 1)



### Cache Coherence Problem (Step 2)



### Cache Coherence Problem (Step 3)



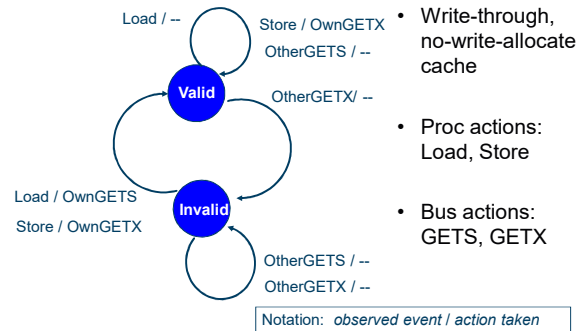
### Snooping Cache-Coherence Protocols

- Each cache controller "snoops" all bus transactions
  - Transaction is relevant if it is for a block this cache contains
  - Take action to ensure coherence
    - Invalidate
    - Update
    - Supply value to requestor if Owner
  - Actions depend on the state of the block and the protocol
- Main memory controller also snoops on bus
  - If no cache is owner, then memory is owner
- Simultaneous operation of independent controllers

## Processor and Bus Actions

- Processor:
  - Load
  - Store
  - Writeback on replacement of modified block
- Bus
  - GetShared (GETS): Get *without* intent to modify, data could come from memory or another cache
  - GetExclusive (GETX): Get *with* intent to modify, must invalidate all other caches' copies
  - PutExclusive (PUTX): cache controller puts contents on bus and memory is updated
  - Definition: **cache-to-cache transfer** occurs when another cache satisfies GETS or GETX request
- Let's draw it!

## Simple 2-State Invalidate Snooping Protocol

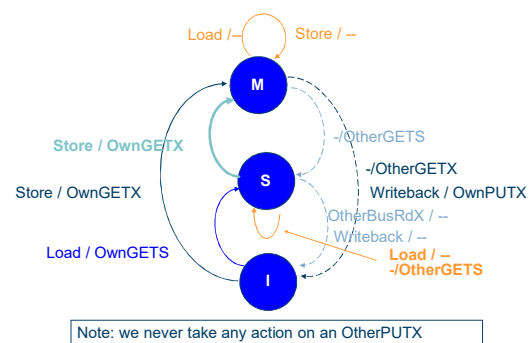


- Write-through, no-write-allocate cache
- Proc actions: Load, Store
- Bus actions: GETS, GETX

## A 3-State Write-Back Invalidation Protocol

- 2-State Protocol
  - + Simple hardware and protocol
  - Uses lots of bandwidth (every write goes on bus!)
- 3-State Protocol (MSI)
  - Modified
    - One cache exclusively has valid (modified) copy → Owner
    - Memory is stale
  - Shared
    - ≥ 1 cache and memory have valid copy (memory = owner)
  - Invalid (only memory has valid copy and memory is owner)
- Must invalidate all other copies before entering Modified state
- Requires bus transaction (order and invalidate)

## MSI State Diagram



## An MSI Protocol Example

Proc Action	P1 State	P2 state	P3 state	Bus Act	Data from
initially	I	I	I		
1. P1 load u	I→S	I	I	GETS	Memory
2. P3 load u	S	I	I→S	GETS	Memory
3. P3 store u	S→I	I	S→M	GETX	Memory or P1 (?)
4. P1 load u	I→S	I	M→S	GETS	P3's cache
5. P2 load u	S	I→S	S	GETS	Memory

- Single writer, multiple reader protocol
- Why Modified to Shared in line 4?
- What if not in any cache? Memory responds
- Read then Write produces 2 bus transactions
  - Slow and wasteful of bandwidth for a common sequence of actions

## Multicore and Multithreaded Processors

- Why multicore?
- Thread-level parallelism
- Multithreaded cores
- Multiprocessors
- Design issues
- Examples

## Some Real-World Multicores

- Intel/AMD 2/4/8-core chips
  - Pretty standard
- Sun's Niagara (UltraSPARC T1-T3)
  - 4-16 simple, in-order, multithreaded cores
- Sun's Rock processor: 16 cores
- Cell Broadband Engine: in PlayStation 3
- Intel's MIC/Larrabee chip: 80 simple x86 cores in a ring
- Cisco CRS-1 Processor: 188 in-order cores
- Graphics processing units (GPUs): hundreds of "cores"